

Homework 04: Fitting, network flow, and project scheduling

ISyE 524 — Fall 2026

Due: Monday, October 5, 2026, at 11:59 p.m. (Madison time).

Complete all three problems. Write each mathematical formulation before implementing it in Julia/JuMP. Define variables, domains, units where applicable, objective, and constraints. Use continuous variables throughout.

This assignment uses the piecewise-linear, minimax, network-flow, and project-scheduling material associated with Lectures 8–9.

When submitting your problem to Gradescope, be sure to appropriately mark which pages of the PDF submission belong to which problem. You won't get full marks unless you do that.

Getting started and written answers

Follow [Assignments and PDF Submission](#). In VS Code, use **Tasks: Run Task** to run **ISyE 524: Update course repository**, then **ISyE 524: Start an assignment from its template**, entering hw04. Open `student-work/hw04/hw04.ipynb`. An existing working copy will not be overwritten.

For **every written part**, type LaTeX mathematics in Markdown or insert a clear photograph/scan of handwritten work. Include the answer and reasoning together and label the subparts; no duplicate transcription is required. Julia/JuMP implementations belong in code cells.

Use $x_1 \geq 0$ for inline mathematics and $x_1 + x_2 \leq 5$ for display mathematics. Run the Markdown cell to render it.

For handwritten work, edit a Markdown cell in VS Code, drag in a cropped PNG/JPEG, and choose **Insert Image as Attachment**. Alternatively, put the image beside your working notebook under `student-work/hw04/images/` and insert `![My answer](images/problem1.png)`. Keep external images with the notebook. See [Images and handwritten work](#).

Back up `student-work/hw04/` separately; Git does not back up your answers or images. After completing the notebook, export it with the course PDF task, inspect `submissions/hw04.pdf` for readable mathematics, tables, outputs, and images, and submit that PDF to Gradescope.

Computational setup

Use JuMP and HiGHS for all models. Print the solver's termination status and check for `MOI.OPTIMAL` and an available solution before reading values. Check constraints using unrounded values.

The small fitting instance, network, and stadium data are supplied below. Keep `data/regression-large.csv` with this notebook; Problem 1(c) reads it with `CSV.jl` and `DataFrames.jl`. All supplied cells run without completed answers. Avoid printing entire large matrices or models in your submission.

AI tools may help debug your own code, explain syntax or error messages, and critique your work. They should not replace constructing your formulation, implementation, or written solution. Complete the required collaboration and LLM-use statement at the end.

```
using JuMP, HiGHS, CSV, DataFrames, Printf
import MathOptInterface as MOI
```

1: Fitting equations when they do not agree

Several measurements of an unknown vector x give the following inconsistent system:

$$\begin{aligned}2x_1 - x_2 + 3x_3 &= 15, \\-x_1 + 4x_2 + 2x_4 &= -9, \\3x_1 + x_2 - 2x_3 + x_4 &= -5, \\2x_2 + x_3 - 3x_4 &= 8, \\x_1 - 2x_2 + 4x_4 &= -3, \\4x_1 + x_3 - x_4 &= 17, \\-2x_1 + 3x_2 + 2x_3 + x_4 &= -1, \\x_1 + x_2 - x_3 + 2x_4 &= -8.\end{aligned}$$

The coefficients x_j are **unrestricted in sign**. We seek an approximate fit, rather than insist that every equation hold exactly.

More generally, let M index measurements and N index unknown coefficients. The data are a_{ij} and b_i . For any proposed x , define its signed error and absolute residual by

$$e_i(x) = \sum_{j \in N} a_{ij}x_j - b_i, \quad r_i(x) = |e_i(x)| \quad (i \in M).$$

1(a): Minimum total absolute error — the L_1 fit

Write a **general indexed LP** that minimizes $\sum_{i \in M} r_i(x)$. A formulation whose objective still contains absolute values is not yet an LP formulation.

Your mathematical formulation and reasoning:

[Type your answer here, or insert a clear photograph/scan of handwritten work.]

$$\min \sum_{i \in M} r_i(x)$$

Subject To:

$$r_i(x) = e_i(x) \quad (i \in M)$$

$$e_i(x) \geq \sum_{j \in N} a_{ij}x_j - b_i$$

$$e_i(x) \geq -\sum_{j \in N} a_{ij}x_j - b_i$$

$$e_i(x) \geq 0$$

Supplied data for the small instance

Here $A_small[i, j]$ is a_{ij} and $b_small[i]$ is b_i . Implement your model from 1(a), solve this instance, and report the four coefficients, minimum total absolute error, and the absolute residual of each equation.

```
A_small = Float64[
    2 -1  3  0
   -1  4  0  2
    3  1 -2  1
    0  2  1 -3
    1 -2  0  4
    4  0  1 -1
   -2  3  2  1
    1  1 -1  2
]
b_small = Float64[15, -9, -5, 8, -3, 17, -1, -8]
println("Small instance: ", size(A_small, 1), " equations, ",
        size(A_small, 2), " unknown coefficients.")

m, n = size(A_small)
```

Small instance: 8 equations, 4 unknown coefficients.

(8, 4)

```
# Write your Julia/JuMP implementation here.
small=Model(HiGHS.Optimizer)

@variable(small, x[j in 1:n])

@variable(small, e[i in 1:m] >= 0)

@constraint(small, E_upper[i in 1:m],
    e[i] >= sum(A_small[i, j]*x[j]-b_small[i] for j in 1:n))

@constraint(small, E_lower[i in 1:m],
    e[i] >= -sum(A_small[i, j]*x[j]-b_small[i] for j in 1:n))

@objective(small, Min, sum(e[i] for i in 1:m))

optimize!(small)
status = termination_status(small)
println("Termination status: ", status)
status == MOI.OPTIMAL || error("HiGHS stopped with status $(status).")
is_solved_and_feasible(small) || error("No feasible optimal solution is available")

minimum_residual = objective_value(small)

println("Minimum Absolute error: ", minimum_residual)
println("")

xvalues = DataFrame(
    x_values = value.(x),
)
```

Running HiGHS 1.15.1 (git hash: 04024d701f): Copyright (c) 2026 under MIT license
Includes third-party software components, see THIRD_PARTY_NOTICES.md for full details
Using BLAS: blastrampoline
LP has 16 rows; 12 cols; 70 nonzeros
Coefficient ranges:
Matrix [1e+00, 4e+00]
Cost [1e+00, 1e+00]
Bound [0e+00, 0e+00]
RHS [4e+00, 7e+01]
Presolving model
16 rows, 12 cols, 70 nonzeros 0s
16 rows, 12 cols, 70 nonzeros 0s
Presolve reductions: rows 16(-0); columns 12(-0); nonzeros 70(-0) - Not reduced
Problem not reduced by presolve: solving the LP
Using dual simplex solver

Iteration	Objective	Infeasibilities	num(sum)
0	0.0000000000e+00	Pr: 8(264); Du: 0(7.61267e-12)	0.0s
11	3.1440000000e+01	Pr: 0(0)	0.0s

Model status : Optimal
 Simplex iterations: 11
 Objective value : 3.1440000000e+01
 P-D objective error : 3.8930801114e-16
 HiGHS run time : 0.00
 Termination status: OPTIMAL
 Minimum Absolute error: 31.44000000000001

	x_values
	Float64
1	6.8
2	-3.56
3	14.28
4	-8.28

```

rvalues = DataFrame(
    residuals = value.(e),
)

```

	residuals
	Float64
1	-0.0
2	1.6
3	0.0
4	0.0
5	7.2
6	18.24
7	0.0
8	4.4

Your results and interpretation:

[Type your answer here, or insert a clear photograph/scan of handwritten work.]

The minimum absolute error is 31.44 with individual residuals of 0, 1.6, 0, 0, 7.2, 18.24, 0, 4.4 for equations 1-8 respectively. The final values for x are $x_1 = 6.8$, $x_2 = -3.56$, $x_3 = 14.28$, $x_4 = -8.28$

1(b): Minimum worst error — the L_∞ fit

Write a general indexed LP that minimizes $\max_{i \in M} r_i(x)$, and explain its equivalence to the minimax problem. Implement it for the same small instance.

Report the four fitted coefficients and every absolute residual. Make a two-row comparison table: for each fit, report **both** the sum of absolute residuals and the largest absolute residual. Explain

what each objective sacrifices to improve its own measure. Would a single signed bound on the errors suffice? Explain.

Your mathematical formulation and reasoning:

[Type your answer here, or insert a clear photograph/scan of handwritten work.]

$$\min \quad t$$

Subject To:

$$\begin{aligned} t &\geq r_i(x) \quad (i \in M) \\ r_i(x) &= e_i(x) \quad (i \in M) \\ e_i(x) &\geq \sum_{j \in N} a_{ij}x_j - b_i \\ e_i(x) &\geq -\sum_{j \in N} a_{ij}x_j - b_i \\ e_i(x) &\geq 0 \end{aligned}$$

To produce a formulation equivalent to the minimax problem I added a variable t which must be greater than or equal to the largest residual value, (this is the maximum part). Then I make a new objective function that minimizes t .

```
# Write your Julia/JuMP implementation here.
inf=Model(HiGHS.Optimizer)

@variable(inf, x[j in 1:n])

@variable(inf, e[i in 1:m] >= 0)

@variable(inf, t >= 0)

@constraint(inf, E_upper[i in 1:m],
    e[i] >= sum(A_small[i, j]*x[j]-b_small[i] for j in 1:n))

@constraint(inf, E_lower[i in 1:m],
    e[i] >= -sum(A_small[i, j]*x[j]-b_small[i] for j in 1:n))

@constraint(inf, ma[i in 1:m],
    t >= e[i])

@objective(inf, Min, t)

optimize!(inf)
status = termination_status(inf)
println("Termination status: ", status)
```

```

status == MOI.OPTIMAL || error("HiGHS stopped with status $(status).")
is_solved_and_feasible(inf) || error("No feasible optimal solution is available.

largest_residual = objective_value(inf)

println("Minimum Absolute error: ", largest_residual)
println("")

xvalues_inf = DataFrame(
    x_values_inf = value.(x),
)

combined_x=hcat(xvalues,xvalues_inf) #used Gemini to find this hcat function

```

Running HiGHS 1.15.1 (git hash: 04024d701f): Copyright (c) 2026 under MIT license
Includes third-party software components, see THIRD_PARTY_NOTICES.md for full details
Using BLAS: blastrampoline
LP has 24 rows; 13 cols; 86 nonzeros

Coefficient ranges:

Matrix	[1e+00, 4e+00]
Cost	[1e+00, 1e+00]
Bound	[0e+00, 0e+00]
RHS	[4e+00, 7e+01]

Presolving model

24 rows, 13 cols, 86 nonzeros 0s

21 rows, 10 cols, 80 nonzeros 0s

18 rows, 7 cols, 74 nonzeros 0s

16 rows, 5 cols, 70 nonzeros 0s

16 rows, 5 cols, 70 nonzeros 0s

Presolve reductions: rows 16(-8); columns 5(-8); nonzeros 70(-16)

Solving the presolved LP

Using dual simplex solver

Iteration	Objective	Infeasibilities	num(sum)
0	0.000000000000e+00	Pr: 8(264); Du: 0(7.24557e-12)	0.0s
7	6.5523809524e+00	Pr: 0(0)	0.0s

Performed postsolve

Solving the original LP from the solution after postsolve

Model status : Optimal

Simplex iterations: 7

Objective value : 6.5523809524e+00

P-D objective error : 6.2970110782e-16

HiGHS run time : 0.00

Termination status: OPTIMAL

Minimum Absolute error: 6.5523809523809575

	x_values	x_values_inf
	Float64	Float64
1	6.8	9.48571
2	-3.56	-3.80952
3	14.28	14.5905
4	-8.28	-8.91429

```
rvalues_inf = DataFrame(
    residuals_inf = value.(e),
)

combined_r=hcat(rvalues, rvalues_inf)
```

	residuals	residuals_inf
	Float64	Float64
1	-0.0	6.55238
2	1.6	6.55238
3	0.0	6.55238
4	0.0	6.55238
5	7.2	6.55238
6	18.24	6.55238
7	0.0	6.55238
8	4.4	6.55238

```
println("Total Error is: ", sum.(eachcol(rvalues_inf)), " for the infinite norm case")
println("Total Error is: ", sum.(eachcol(rvalues)), " for the first norm case")
println("")
println("Largest infinite norm residual is: ", largest_residual)
println("Largest first norm residual is: ", maximum(Matrix(rvalues)))
```

Total Error is: [52.41904761904766] for the infinite norm case
Total Error is: [31.440000000000001] for the first norm case

Largest infinite norm residual is: 6.5523809523809575
Largest first norm residual is: 18.240000000000002

Your results and interpretation:

[Type your answer here, or insert a clear photograph/scan of handwritten work.]

In the first norm residual you compromise your highest residual for a lower total residual, thus meaning you fit most of the equations really good but dont fit one partilarly well. On the other hand in the inifinte norm case you sacrifise overall error for having very low error to each equation. This means you fit each eqation relitively equally but dont fit the whole system as well as the first norm case.

Single signed bounds on the error function in a error minimization problem would not suffice as it would cause an unbounded problem. This is because the error values being infinite and negative would create the “smallest” error possible even though the error would be huge and just negative. This is why error minization must always be absolute valued.

1(c): The same models at a larger scale

The file `data/regression-large.csv` contains a **fixed synthetic instance with 1,200 equations and 60 unknown coefficients**. All entries are dimensionless. Each row contains:

- equation: a row identifier, not a coefficient;
- a01 through a60: the coefficients $a_{i1}, \dots, a_{i60}$;
- b: the observed right-hand side b_i .

Run the supplied CSV/DataFrames cell and inspect the preview. Reuse your general formulations from 1(a)–1(b) to solve **both** fits. Report:

1. A two-row table of total absolute error and maximum absolute error, computed from the fitted coefficients and the original data.

No new mathematical formulation or printout of all 1,200 residuals is required.

```
# Locate this assignment when the kernel starts beside the notebook or
# at the repository root. Prefer an existing personal copy.
assignment_dir = if isfile(joinpath(pwd(), "hw04.ipynb"))
    pwd()
elseif isdir(joinpath(pwd(), "student-work", "hw04"))
    joinpath(pwd(), "student-work", "hw04")
else
    joinpath(pwd(), "assignments", "hw04")
end
data_dir = joinpath(assignment_dir, "data")
data_path = joinpath(data_dir, "regression-large.csv")
isfile(data_path) || error(
    "Missing regression-large.csv in $(data_dir). " *
    "Keep the complete data folder beside your HW04 notebook."
)
regression_data = CSV.read(data_path, DataFrame)
coefficient_columns = filter(
    name -> startswith(name, "a"), names(regression_data))
A_large = Matrix{Float64}(regression_data[:, coefficient_columns])
b_large = Vector{Float64}(regression_data.b)
equation_ids = regression_data.equation
M_large = axes(A_large, 1)
N_large = axes(A_large, 2)

@assert length(unique(equation_ids)) == length(equation_ids)
@assert size(A_large, 1) == length(b_large)
@assert all(isfinite, A_large) && all(isfinite, b_large)
println("CSV instance: ", length(M_large), " equations, ",
        length(N_large), " unknown coefficients.")
first(regression_data[:, [:equation, :a01, :a02, :a03, :b]], 5)
```

Base.OneTo(60)

CSV instance: 1200 equations, 60 unknown coefficients.

	equation	a01	a02	a03	b
	Int64	Int64	Int64	Int64	Float64
1	1	-3	1	-1	76.3728
2	2	5	0	-1	42.5359
3	3	-3	-5	5	37.3637
4	4	3	5	3	-12.6925
5	5	4	-2	-3	-36.9703

```
# Write your Julia/JuMP implementation here.

n_large=length(N_large)
m_large=length(M_large)

large=Model(HiGHS.Optimizer)

@variable(large, x[j in 1:n_large])

@variable(large, e[i in 1:m_large] >= 0)

@constraint(large, E_upper[i in 1:m_large],
    e[i] >= sum(A_large[i, j]*x[j]-b_large[i] for j in 1:n_large))

@constraint(large, E_lower[i in 1:m_large],
    e[i] >= -sum(A_large[i, j]*x[j]-b_large[i] for j in 1:n_large))

@objective(large, Min, sum(e[i] for i in 1:m_large))

optimize!(large)
status = termination_status(large)
println("Termination status: ", status)
status == MOI.OPTIMAL || error("HiGHS stopped with status $(status).")
is_solved_and_feasible(large) || error("No feasible optimal solution is available")

minimum_residual_large = objective_value(large)

println("Minimum Absolute error: ", minimum_residual_large)
println("")

xvalues_large = DataFrame(
    x_values_large = value.(x),
)

rvalues_large = DataFrame(
    residuals_large = value.(e),
)
```

```

Running HiGHS 1.15.1 (git hash: 04024d701f): Copyright (c) 2026 under MIT license
Includes third-party software components, see THIRD_PARTY_NOTICES.md for full details
Using BLAS: blastrampoline
LP has 2400 rows; 1260 cols; 133388 nonzeros
Coefficient ranges:
  Matrix  [1e+00, 5e+00]
  Cost     [1e+00, 1e+00]
  Bound    [0e+00, 0e+00]
  RHS      [9e+00, 1e+04]
Presolving model
2400 rows, 1260 cols, 133388 nonzeros 0s
2400 rows, 1260 cols, 133388 nonzeros 0s
Presolve reductions: rows 2400(-0); columns 1260(-0); nonzeros 133388(-0) - Not reduced
Problem not reduced by presolve: solving the LP
Using dual simplex solver
  Iteration      Objective      Infeasibilities num(sum)
           0      0.0000000000e+00 Pr: 1200(3.14612e+06); Du: 0(1.89107e-09) 0.0s
        2272      8.2851611436e+04 Pr: 0(0) 0.8s

Model status      : Optimal
Simplex iterations: 2272
Objective value    : 8.2851611436e+04
P-D objective error: 1.5807350565e-15
HiGHS run time     : 0.79
Termination status: OPTIMAL
Minimum Absolute error: 82851.61143579037

```

	residuals_large
	Float64
1	15.4508
2	8.325
3	8.01585
4	34.6843
5	11.6383
6	42.597
7	28.2575
8	5.85821
9	38.3587
10	48.2598
11	25.784
12	1.38875
13	47.0539
14	122.26
15	69.683
16	2.41719e-8
17	50.9782
18	57.6131
19	96.6839
20	58.7585
21	12.6313
22	44.1619
23	36.3583
24	10.0161
...	...

```

inf_large=Model(HiGHS.Optimizer)

@variable(inf_large, x[j in 1:n_large])

@variable(inf_large, e[i in 1:m_large] >= 0)

@variable(inf_large, t >= 0)

@constraint(inf_large, E_upper[i in 1:m_large],
    e[i] >= sum(A_large[i, j]*x[j]-b_large[i] for j in 1:n_large))

@constraint(inf_large, E_lower[i in 1:m_large],
    e[i] >= -sum(A_large[i, j]*x[j]-b_large[i] for j in 1:n_large))

@constraint(inf_large, ma[i in 1:m_large],
    t >= e[i])

@objective(inf_large, Min, t)

optimize!(inf_large)

```

```

status = termination_status(inf_large)
println("Termination status: ", status)
status == MOI.OPTIMAL || error("HiGHS stopped with status $(status).")
is_solved_and_feasible(inf) || error("No feasible optimal solution is available.

largest_residual_large = objective_value(inf_large)

println("Minimum Absolute error: ", largest_residual_large)
println("")

xvalues_inf_large = DataFrame(
    x_values_inf_large = value.(x),
)

combined_x_large=hcat(xvalues_large,xvalues_inf_large) #used Gemini to find this

```

```

Running HiGHS 1.15.1 (git hash: 04024d701f): Copyright (c) 2026 under MIT license
Includes third-party software components, see THIRD_PARTY_NOTICES.md for full details
Using BLAS: blastrampoline
LP has 3600 rows; 1261 cols; 135788 nonzeros
Coefficient ranges:
  Matrix  [1e+00, 5e+00]
  Cost     [1e+00, 1e+00]
  Bound    [0e+00, 0e+00]
  RHS      [9e+00, 1e+04]
Presolving model
3600 rows, 1261 cols, 135788 nonzeros 0s
3235 rows, 896 cols, 135058 nonzeros 2s
2925 rows, 586 cols, 134438 nonzeros 6s
2661 rows, 322 cols, 133910 nonzeros 10s
2437 rows, 98 cols, 133462 nonzeros 14s
2400 rows, 61 cols, 133388 nonzeros 15s
Presolve reductions: rows 2400(-1200); columns 61(-1200); nonzeros 133388(-
2400)
Solving the presolved LP
Using dual simplex solver
  Iteration      Objective      Infeasibilities num(sum)
          0          0.0000000000e+00 Pr: 1200(3.14612e+06); Du: 0(2.12961e-
09) 15.2s
        229          9.7696438091e+02 Pr: 0(0) 15.2s

Performed postsolve
Solving the original LP from the solution after postsolve

Model status      : Optimal
Simplex iterations: 229

```

Objective value : 9.7696438091e+02
 P-D objective error : 4.7104702920e-15
 HiGHS run time : 15.27
 Termination status: OPTIMAL
 Minimum Absolute error: 976.9643809101464

	x_values_large	x_values_inf_large
	Float64	Float64
1	-60.4496	-42.3492
2	90.7319	102.142
3	-120.082	-103.992
4	150.56	164.138
5	-181.514	-124.107
6	210.006	225.164
7	-29.7231	-18.6531
8	59.8513	64.6369
9	-89.1353	-75.2902
10	120.14	96.3008
11	-149.302	-114.403
12	179.506	161.621
13	-210.478	-199.659
14	31.0793	34.8112
15	-60.0259	-65.3154
16	89.1952	74.2962
17	-119.329	-112.057
18	149.406	134.994
19	-179.589	-157.121
20	209.928	188.796
21	-30.2983	-56.2832
22	61.0331	77.085
23	-88.9987	-79.1117
24	119.946	138.969
...	...	...

```

rvalues_inf_large = DataFrame(
    residuals_inf_large = value.(e),
)

combined_r_large=hcat(rvalues_large, rvalues_inf_large)

```

	residuals_large	residuals_inf_large
	Float64	Float64
1	15.4508	976.964
2	8.325	976.964
3	8.01585	976.964
4	34.6843	976.964
5	11.6383	976.964
6	42.597	976.964
7	28.2575	976.964
8	5.85821	976.964
9	38.3587	976.964
10	48.2598	976.964
11	25.784	976.964
12	1.38875	976.964
13	47.0539	976.964
14	122.26	976.964
15	69.683	976.964
16	2.41719e-8	976.964
17	50.9782	976.964
18	57.6131	976.964
19	96.6839	976.964
20	58.7585	976.964
21	12.6313	976.964
22	44.1619	976.964
23	36.3583	976.964
24	10.0161	976.964
...	...	...

```
println("Total Error is: ", sum.(eachcol(rvalues_inf_large)), " for the infinite
println("Total Error is: ", sum.(eachcol(rvalues_large)), " for the first norm ca
println("")
println("Largest infinite norm residual is: ", largest_residual_large)
println("Largest first norm residual is: ", maximum(Matrix(rvalues_large)))
```

Total Error is: [1.1723572570921765e6] for the infinite norm case

Total Error is: [82851.61143579039] for the first norm case

Largest infinite norm residual is: 976.9643809101464

Largest first norm residual is: 1555.6710856813586

Your results and interpretation:

[Type your answer here, or insert a clear photograph/scan of handwritten work.]

Norm Type	Total Absolute Error	Maximum Absolute error
1	82851.6	1555.7
∞	1.1723×10^6	977.0

Norm Type	Total Absolute Error	Maximum Absolute error
-----------	----------------------	------------------------

Here the absolute error explodes when you impose the minimax on the equation because there is so many equations.

2: Minimum-cost distribution through a network

A distributor must move one interchangeable product through ten locations. Locations 1, 2, and 9 supply product; locations 6, 7, and 8 require it. The other locations are transfer points. Product can also pass through a location that has supply or demand.

The planning period has the following **net supplies** b_i : a positive number means supply, a negative number means demand, and zero means pure transshipment.

Node i	1	2	3	4	5	6	7	8	9	10
b_i (tons)	10	15	0	0	0	-12	-11	-7	5	0

Only the directed arcs in the following table are available. An arc from i to j does **not** imply that the reverse arc exists. Every arc has lower bound zero, cost c_{ij} dollars per ton, and capacity u_{ij} tons. All listed supply must be shipped and all listed demand must be met exactly. There is no loss in transit.

From	To	Cost (\$/ton)	Capacity (tons)
1	3	5	15
2	4	4	20
3	4	2	12
3	5	6	10
3	6	5	10
4	5	1	10
4	8	2	10
5	3	4	8
5	6	6	10
5	7	3	15
8	7	4	8
9	3	3	5
9	10	2	5
10	5	1	6
10	8	4	4
3	10	2	6

2(a): Mathematical formulation

Write a general minimum-cost network-flow LP using node set N , arc set A , net supplies b_i , costs c_{ij} , and capacities u_{ij} . Let f_{ij} be the flow on arc (i, j) . Explain your flow-balance sign convention and why the sum of all net supplies must be zero. Use variables only for existing arcs.

Your mathematical formulation and reasoning:

[Type your answer here, or insert a clear photograph/scan of handwritten work.]

$$\min \sum_{(i,j) \in \mathcal{E}} f_{ij} c_{ij}$$

Subject To:

$$\begin{aligned} \sum_{(i,j) \in \mathcal{E}} f_{ij} - \sum_{(i,j) \in \mathcal{E}} f_{ji} &= b_i \quad \forall i \in N \\ 0 &\leq f_{ij} \leq c_{ij} \quad \forall (i, j) \in \mathcal{E} \end{aligned}$$

Here it is important to have the sign convention be positive for flow out and negative for flow in so that it agrees with the sign convention of the problem where demand is negative and supply is positive. Along with this it is important that the sum of all net supplies is zero so that this problem is balanced. If this problem were unbalanced and had exact node balance equalities like is used in this problem it would be infeasible.

2(b): Implement, solve, and check

Implement your formulation using the supplied data. Report the minimum total cost and all positive arc flows. Give a check of the net outflow at every node and verify all capacity limits. Identify the arcs operating at capacity. Interpret the result as a shipment plan; integer declarations are not required.

```
N = 1:10
b = Dict{zip(N, [10, 15, 0, 0, 0, -12, -11, -7, 5, 0])}
# Each row is (from, to, cost, capacity).
arc_data = [
    (1, 3, 5, 15), (2, 4, 4, 20), (3, 4, 2, 12),
    (3, 5, 6, 10), (3, 6, 5, 10), (4, 5, 1, 10),
    (4, 8, 2, 10), (5, 3, 4, 8), (5, 6, 6, 10),
    (5, 7, 3, 15), (8, 7, 4, 8), (9, 3, 3, 5),
    (9, 10, 2, 5), (10, 5, 1, 6), (10, 8, 4, 4),
    (3, 10, 2, 6)
]
A = [(i, j) for (i, j, cost, capacity) in arc_data]
c = Dict{((i, j) => cost) for (i, j, cost, capacity) in arc_data}
u = Dict{((i, j) => capacity) for (i, j, cost, capacity) in arc_data}
println("Network: ", length(N), " nodes, ",
```

```
length(A), " directed arcs.")
```

Network: 10 nodes, 16 directed arcs.

```
# Write your Julia/JuMP implementation here.

mcn=Model(HiGHS.Optimizer)

@variable(mcn, f[(i,j) in A] >= 0)

@constraint(mcn, flowbalance[i in N],
    sum(f[(i,j)] for j in N if (i,j) in A) -
    sum(f[(j,i)] for j in N if (j,i) in A)
    == b[i])

@constraint(mcn, capacity[(i,j) in A],
    (f[(i,j)] <= u[(i,j)]))

@objective(mcn, Min,
    sum(c[i,j]*f[(i,j)] for (i,j) in A))

optimize!(mcn)
status = termination_status(mcn)
println("Termination status: ", status)
status == MOI.OPTIMAL || error("HiGHS stopped with status $(status).")
is_solved_and_feasible(mcn) || error("No feasible optimal solution is available.")

min_cost= objective_value(mcn)

println("Minimum Cost if: ", min_cost, " Dollars")
println()

flows = DataFrame(
    flows = value.(f[(i,j)] for (i,j) in A),
    to_from = collect(A),
    capacity = collect(u[(i,j)] for (i,j) in A)
)
```

Running HiGHS 1.15.1 (git hash: 04024d701f): Copyright (c) 2026 under MIT license
Includes third-party software components, see THIRD_PARTY_NOTICES.md for full details
Using BLAS: blastrampoline

LP has 26 rows; 16 cols; 48 nonzeros

Coefficient ranges:

Matrix [1e+00, 1e+00]

Cost [1e+00, 6e+00]

Bound [0e+00, 0e+00]

RHS [4e+00, 2e+01]

Presolving model

5 rows, 11 cols, 22 nonzeros 0s

Dependent equations search running on 5 equations with time limit of 1000.00s

Dependent equations search removed 1 rows and 6 nonzeros in 0.00s (limit = 1000.

4 rows, 11 cols, 16 nonzeros 0s

Presolve reductions: rows 4(-22); columns 11(-5); nonzeros 16(-32)

Solving the presolved LP

Using dual simplex solver

Iteration	Objective	Infeasibilities num(sum)
0	0.0000000000e+00	Ph1: 0(0) 0.0s
3	2.4200000000e+02	Pr: 0(0) 0.0s

Performed postsolve

Solving the original LP from the solution after postsolve

Model status : Optimal

Simplex iterations: 3

Objective value : 2.4200000000e+02

P-D objective error : 0.0000000000e+00

HiGHS run time : 0.00

Termination status: OPTIMAL

Minimum Cost if: 242.0 Dollars

	flows	to_from	capacity
	Float64	Tuple...	Int64
1	10.0	(1, 3)	15
2	15.0	(2, 4)	20
3	0.0	(3, 4)	12
4	0.0	(3, 5)	10
5	10.0	(3, 6)	10
6	8.0	(4, 5)	10
7	7.0	(4, 8)	10
8	0.0	(5, 3)	8
9	2.0	(5, 6)	10
10	11.0	(5, 7)	15
11	0.0	(8, 7)	8
12	0.0	(9, 3)	5
13	5.0	(9, 10)	5
14	5.0	(10, 5)	6
15	0.0	(10, 8)	4
16	0.0	(3, 10)	6

```

net_flows = DataFrame(
    net_flow_out = [sum(value(f[(i,j)])) for j in N if (i,j) in A; init=0.0) for i in N
    net_flow_in  = [sum(value(f[(j,i)])) for j in N if (j,i) in A; init=0.0) for i in N
    b_expected=[b[i] for i in N]
)

net_flows.net_flow = net_flows.net_flow_out - net_flows.net_flow_in

net_flows

```

	net_flow_out	net_flow_in	b_expected	net_flow
	Float64	Float64	Int64	Float64
1	10.0	0.0	10	10.0
2	15.0	0.0	15	15.0
3	10.0	10.0	0	0.0
4	15.0	15.0	0	0.0
5	13.0	13.0	0	0.0
6	0.0	12.0	-12	-12.0
7	0.0	11.0	-11	-11.0
8	0.0	7.0	-7	-7.0
9	5.0	0.0	5	5.0
10	5.0	5.0	0	0.0

Your results and interpretation:

[Type your answer here, or insert a clear photograph/scan of handwritten work.]

Implement your formulation using the supplied data. Report the minimum total cost and all positive arc flows. Give a check of the net outflow at every node and verify all capacity limits. Identify the arcs operating at capacity. Interpret the result as a shipment plan; integer declarations are not required.

The minimum cost for this shipping plan is 242 dollars. Shipment Plan shown in image below.




Figure 1: image.png

3: Stadium building

A town council wishes to construct a small stadium to improve local services. A construction company has won the contract and wants to plan the work. Some activities can begin only after their predecessors finish. Activities may otherwise proceed simultaneously; assume sufficient crews and equipment, with no shared-resource restrictions.

Time is measured in **weeks from the start of construction**, and week 0 is the earliest start. The following are the normal activity durations and immediate predecessors.

Task	Description	Duration d_i	Predecessors
1	Install construction site	2	none
2	Terracing	16	1
3	Construct foundations	9	2
4	Access roads and other networks	8	2
5	Erect basement	10	3
6	Main floor	6	4, 5
7	Divide changing rooms	2	4

Task	Description	Duration d_i	Predecessors
8	Electrify terraces	2	6
9	Construct roof	9	4, 6
10	Stadium lighting	5	4
11	Install terraces	3	6
12	Seal roof	2	9
13	Finish changing rooms	1	7
14	Construct ticket office	7	2
15	Secondary access roads	4	4, 14
16	Install signs	3	8, 11, 14
17	Lawn and sports accessories	9	12
18	Hand over building	1	17

The project is complete only when all eighteen tasks are finished. In particular, a task without successors still has to finish. Task 18 retains the predecessors listed in the table.

Let $T = \{1, \dots, 18\}$ and $P = \{(i, j) : i \text{ immediately precedes } j\}$. This problem uses **predecessor first** in each pair. Let x_i denote the start time of task i .

3(a): Earliest completion

Write a general LP for the minimum project completion time. Implement and solve it using the supplied data. Report the minimum time, a feasible start/finish schedule, and a critical path. Explain how the durations along that path certify that no shorter schedule is possible.

Your mathematical formulation and reasoning:

[Type your answer here, or insert a clear photograph/scan of handwritten work.]

$$\min \sum_{(i) \in \mathcal{E}} x_T + d_T$$

Subject To:

$$t_1 = 0$$

$$x_i \geq x_j + x_j \quad \forall j \in \mathcal{P}_i \quad \forall i \in \mathcal{T}$$

```

T = 1:18
d = [2, 16, 9, 8, 10, 6, 2, 2, 9, 5, 3, 2, 1, 7, 4, 3, 9, 1]
pred = [Int[], [1], [2], [2], [3], [4,5], [4], [6],
        [4,6], [4], [6], [9], [7], [2], [4,14], [8,11,14],
        [12], [17]]
P = [(i, j) for j in T for i in pred[j]]
println("Stadium: ", length(T), " tasks and ",
        length(P), " precedence pairs.")

```

Stadium: 18 tasks and 22 precedence pairs.

```
# Write your Julia/JuMP implementation here.

stadium=Model(HiGHS.Optimizer)

@variable(stadium, x[i in T])

@constraint(stadium, start[(i,j) in P], x[j]>=x[i]+d[i])

@constraint(stadium, x[1] == 0)

@objective(stadium, Min, x[length(T)]+d[length(T)])

optimize!(stadium)
status = termination_status(stadium)
println("Termination status: ", status)
status == MOI.OPTIMAL || error("HiGHS stopped with status $(status).")
is_solved_and_feasible(stadium) || error("No feasible optimal solution is available")

shortest_time= objective_value(stadium)

println("Shortest Possible time is: ",shortest_time, " weeks")

println("Precedence-constraint slack (not task float):")
for i in T
    for j in pred[i]
        slack = value(x[i]) - value(x[j]) - d[j];
        if (slack <= .00001)
            println("task ", i, " predecessor ",j, " slack is : ", slack, " ***A")
        else
            println("task ", i, " predecessor ",j, " slack is : ", slack, " ***F")
        end
    end
end

# Forward pass: earliest starts and a dependency order.
earliest = Dict{Int64,Float64}()
order = Int[]
while length(order) < length(T)
    ready = [i for i in T if !haskey(earliest, i) &&
            all(haskey(earliest, j) for j in pred[i])]
    isempty(ready) && error("The precedence graph contains a cycle")
    for i in ready
        earliest[i] = isempty(pred[i]) ? 0.0 :
            maximum(earliest[j] + d[j] for j in pred[i])
    end
end
```

```

        push!(order, i)
    end
end

# Backward pass: latest starts consistent with the minimum makespan.
latest = Dict{Int64,Float64}()
for i in reverse(order)
    successors = [j for j in T if i in pred[j]]
    latest[i] = isempty(successors) ? objective_value(stadium) - d[i] :
        minimum(latest[j] - d[i] for j in successors)
end

total_float = Dict{i => latest[i] - earliest[i] for i in T}
critical = [i for i in order if abs(total_float[i]) <= 1e-7]
println("Critical activities (including dummy start/finish): ", critical)

```

Running HiGHS 1.15.1 (git hash: 04024d701f): Copyright (c) 2026 under MIT license
Includes third-party software components, see THIRD_PARTY_NOTICES.md for full details
Using BLAS: blastrampoline

LP has 23 rows; 18 cols; 45 nonzeros

Coefficient ranges:

Matrix	[1e+00, 1e+00]
Cost	[1e+00, 1e+00]
Bound	[0e+00, 0e+00]
RHS	[2e+00, 2e+01]

Presolving model

5 rows, 6 cols, 10 nonzeros 0s

0 rows, 0 cols, 0 nonzeros 0s

Presolve reductions: rows 0(-23); columns 0(-18); nonzeros 0(-45) - Reduced to empty

Performed postsolve

Solving the original LP from the solution after postsolve

Model status : Optimal
Objective value : 6.4000000000e+01
P-D objective error : 0.0000000000e+00
HiGHS run time : 0.00

Termination status: OPTIMAL

Shortest Possible time is: 64.0 weeks

Precedence-constraint slack (not task float):

task 2 predecessor 1	slack is : 0.0 ***ACTIVE
task 3 predecessor 2	slack is : 0.0 ***ACTIVE
task 4 predecessor 2	slack is : 0.0 ***ACTIVE
task 5 predecessor 3	slack is : 0.0 ***ACTIVE
task 6 predecessor 4	slack is : 11.0 ***FLOAT
task 6 predecessor 5	slack is : 0.0 ***ACTIVE

```

task 7 predecessor 4  slack is : 0.0 ***ACTIVE
task 8 predecessor 6  slack is : 0.0 ***ACTIVE
task 9 predecessor 4  slack is : 17.0 ***FLOAT
task 9 predecessor 6  slack is : 0.0 ***ACTIVE
task 10 predecessor 4  slack is : 0.0 ***ACTIVE
task 11 predecessor 6  slack is : 0.0 ***ACTIVE
task 12 predecessor 9  slack is : 0.0 ***ACTIVE
task 13 predecessor 7  slack is : 0.0 ***ACTIVE
task 14 predecessor 2  slack is : 0.0 ***ACTIVE
task 15 predecessor 4  slack is : 0.0 ***ACTIVE
task 15 predecessor 14  slack is : 1.0 ***FLOAT
task 16 predecessor 8  slack is : 1.0 ***FLOAT
task 16 predecessor 11  slack is : 0.0 ***ACTIVE
task 16 predecessor 14  slack is : 21.0 ***FLOAT
task 17 predecessor 12  slack is : 0.0 ***ACTIVE
task 18 predecessor 17  slack is : 0.0 ***ACTIVE
Critical activities (including dummy start/finish): [1, 2, 3, 5, 6, 9, 12, 17, 18]

```

```

start_time = DataFrame(
    start = value.(x).data,
    duration = collect(d)
)

```

	start	duration
	Float64	Int64
1	-0.0	2
2	2.0	16
3	18.0	9
4	18.0	8
5	27.0	10
6	37.0	6
7	26.0	2
8	43.0	2
9	43.0	9
10	26.0	5
11	43.0	3
12	52.0	2
13	28.0	1
14	18.0	7
15	26.0	4
16	46.0	3
17	54.0	9
18	63.0	1

Your results and interpretation:

[Type your answer here, or insert a clear photograph/scan of handwritten work.]

The minimum time for this project is 64 weeks. This is constrained by a critical path of tasks 1,2,3,5,6,9,12,17,18. This is the critical path because in each case the possible tasks that may follow take less time. Thus if you sum the durations of the critical path you should get the minimum project time.

$$d_1 + d_2 + d_3 + d_5 + d_5 + d_9 + d_{12} + d_{17} + d_{18} = 64, \quad 2 + 16 + 9 + 10 + 6 + 9 + 2 + 9 + 1 = 64$$

Collaboration and LLM-use statement — required

Complete this statement as the **very last part of your homework**. Be specific enough that a reader can understand what work you did yourself and how others or LLMs helped you learn.

- **People:** Give the name of every person you collaborated with, identify the problem/subpart you discussed, and explain the nature of the collaboration. If you worked alone, explicitly state that you did not collaborate with anyone.
- **LLMs/tools:** Name every LLM or AI assistant you used and identify the model/version as precisely as available. If the interface did not show a model/version, say so. If you used none, explicitly state that you did not use an LLM or AI assistant.
- **Specific assistance and learning:** For each tool, identify the problem/subpart, describe what you asked it to help with, what suggestions you used or rejected, and how the interaction helped you understand the mathematics or implementation. Explain how you checked the suggestions. “I used ChatGPT for help” is not sufficiently specific.

Disclosure does not replace the requirement that your submitted formulation, implementation, and explanation be your own work and that you understand them.

People and collaboration:

[Names, problem/subparts, and what you discussed; or explicitly state that you worked alone.]

I worked alone.

LLMs/tools and models:

[Tool names and models/versions, or explicitly state that you used none.]

Google Gemini 3.5 flash Lite

How I used the assistance, what I learned, and how I checked it:

[Specific details for each tool and relevant problem/subpart; state not applicable if you used none.]

I used LLM to find the heat function. I knew how to add dictionaries in python so I prompted this LLM with the way I would do it in python, `concat([df1,df2], axis=1)` and what the equivalent was in Julia. I then implemented it and printed the results to make sure that it was doing what I intended.